

```
!pip install shap
```

```
⇒ Requirement already satisfied: shap in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/di
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.10/di
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/d
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.10/d
Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dis
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pyt
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/pytho
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p
```

```
!pip install liac-arff
```

```
⇒ Requirement already satisfied: liac-arff in /usr/local/lib/python3.10/dist-
```

```
# Install ucimlrepo for fetching the chronic kidney disease dataset
```

```
!pip install ucimlrepo
```

```
⇒ Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-
Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3
Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/d
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pyt
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p
```

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix
import arff
import shap # For Explainable AI
from scipy.io import arff as scipy_arff # To read ARFF files
from io import StringIO # Import StringIO from the io module
from ucimlrepo import fetch_ucirepo
import gym
from gym import spaces
import random
from sklearn.impute import SimpleImputer

```

```

# Fetch Chronic Kidney Disease Dataset
chronic_kidney_disease = fetch_ucirepo(id=336)

```

```

# Features and targets
X_ckd = chronic_kidney_disease.data.features
y_ckd = chronic_kidney_disease.data.targets

```

```

# Display metadata and variable information
print(chronic_kidney_disease.metadata)
print(chronic_kidney_disease.variables)

```

➞ /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: Deprecat and should\_run\_async(code)

```

{'uci_id': 336, 'name': 'Chronic Kidney Disease', 'repository_url': 'https:

```

	name	role	type	demographic	description \
0	age	Feature	Integer	Age	None
1	bp	Feature	Integer	None	blood pressure
2	sg	Feature	Categorical	None	specific gravity
3	al	Feature	Categorical	None	albumin
4	su	Feature	Categorical	None	sugar
5	rbc	Feature	Binary	None	red blood cells
6	pc	Feature	Binary	None	pus cell
7	pcc	Feature	Binary	None	pus cell clumps
8	ba	Feature	Binary	None	bacteria
9	bgr	Feature	Integer	None	blood glucose random
10	bu	Feature	Integer	None	blood urea
11	sc	Feature	Continuous	None	serum creatinine
12	sod	Feature	Integer	None	sodium
13	pot	Feature	Continuous	None	potassium
14	hemo	Feature	Continuous	None	hemoglobin
15	pcv	Feature	Integer	None	packed cell volume
16	wbcc	Feature	Integer	None	white blood cell count

17	rbcc	Feature	Continuous	None	red blood cell count
18	htn	Feature	Binary	None	hypertension
19	dm	Feature	Binary	None	diabetes mellitus
20	cad	Feature	Binary	None	coronary artery disease
21	appet	Feature	Binary	None	appetite
22	pe	Feature	Binary	None	pedal edema
23	ane	Feature	Binary	None	anemia
24	class	Target	Binary	None	ckd or not ckd

	units	missing_values
0	year	yes
1	mm/Hg	yes
2	None	yes
3	None	yes
4	None	yes
5	None	yes
6	None	yes
7	None	yes
8	None	yes
9	mgs/dl	yes
10	mgs/dl	yes
11	mgs/dl	yes
12	mEq/L	yes
13	mEq/L	yes
14	gms	yes
15	None	yes
16	cells/cmm	yes
17	millions/cmm	yes
18	None	yes
19	None	yes
20	None	yes
21	None	yes
22	None	yes
23	None	yes
24	None	no

```
from ucimlrepo import fetch_ucirepo
```

```
heart_disease = fetch_ucirepo(id=45)
```

```
# data (as pandas dataframes)
X = heart_disease.data.features
y = heart_disease.data.targets
print(heart_disease.metadata)
```

```
print(heart_disease.variables)
```

```
➞ /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: Deprecat
and should_run_async(code)
```

```
{'uci_id': 45, 'name': 'Heart Disease', 'repository_url': 'https://archive.
```

	name	role	type	demographic \
0	age	Feature	Integer	Age
1	sex	Feature	Categorical	Sex
2	cp	Feature	Categorical	None
3	trestbps	Feature	Integer	None
4	chol	Feature	Integer	None
5	fbs	Feature	Categorical	None
6	restecg	Feature	Categorical	None
7	thalach	Feature	Integer	None
8	exang	Feature	Categorical	None
9	oldpeak	Feature	Integer	None
10	slope	Feature	Categorical	None
11	ca	Feature	Integer	None
12	thal	Feature	Categorical	None
13	num	Target	Integer	None

		description	units	missing_values
0		None	years	no
1		None	None	no
2		None	None	no
3	resting blood pressure (on admission to the ho...		mm Hg	no
4		serum cholestoral	mg/dl	no
5		fasting blood sugar > 120 mg/dl	None	no
6		None	None	no
7		maximum heart rate achieved	None	no
8		exercise induced angina	None	no
9	ST depression induced by exercise relative to ...		None	no
10		None	None	no
11	number of major vessels (0-3) colored by flour...		None	yes
12		None	None	yes
13		diagnosis of heart disease	None	no

```
# Load other datasets
```

```
bpx_data = pd.read_sas('BPX_J.XPT') # Blood pressure data
```

```
demo_data = pd.read_sas('DEMO_J.XPT') # Demographics data
```

```
diq_data = pd.read_sas('DIQ_J.XPT') # Diabetes-related information
```

```
dr1tot_data = pd.read_sas('DR1TOT_J.XPT') # Dietary data
glu_data = pd.read_sas('GLU_J.XPT') # Glucose data

# Load diabetic data
diabetic_data = pd.read_csv('diabetic_data.csv')
```

```
# Display first few rows of datasets
print(bpx_data.head())
print(demo_data.head())
print(diag_data.head())
print(dr1tot_data.head())
print(glu_data.head())
print(diabetic_data.head())
```

→ /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: Deprecat  
and should\_run\_async(code)

	SEQN	PEASCCT1	BPXCHR	BPAARM	BPACSZ	BPXPLS	BPXPULS	BPXPTY	BPXM
0	93703.0	NaN	120.0	NaN	NaN	NaN	1.0	NaN	N
1	93704.0	NaN	114.0	NaN	NaN	NaN	1.0	NaN	N
2	93705.0	NaN	NaN	1.0	4.0	52.0	1.0	1.0	220
3	93706.0	NaN	NaN	1.0	3.0	82.0	1.0	1.0	140
4	93707.0	NaN	NaN	1.0	2.0	100.0	1.0	1.0	140

	BPXSY1	...	BPAEN1	BPXSY2	BPXDI2	BPAEN2	BPXSY3	BPXDI3	BPAEN3	\
0	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
1	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
2	NaN	...	NaN	NaN	NaN	NaN	202.0	62.0	2.0	
3	112.0	...	2.0	114.0	70.0	2.0	108.0	76.0	2.0	
4	128.0	...	2.0	128.0	46.0	2.0	128.0	58.0	2.0	

	BPXSY4	BPXDI4	BPAEN4
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	198.0	74.0	2.0
3	NaN	NaN	NaN
4	NaN	NaN	NaN

[5 rows x 21 columns]

	SEQN	SDDSRVYR	RIDSTATR	RIAGENDR	RIDAGEYR	RIDAGEMN	RIDRETH1	\
0	93703.0	10.0	2.0	2.0	2.0	NaN	5.0	
1	93704.0	10.0	2.0	1.0	2.0	NaN	3.0	
2	93705.0	10.0	2.0	2.0	66.0	NaN	4.0	
3	93706.0	10.0	2.0	1.0	18.0	NaN	5.0	
4	93707.0	10.0	2.0	1.0	13.0	NaN	5.0	

	RIDRETH3	RIDEXMON	RIDEXAGM	...	DMDHREDZ	DMDHRMAZ	DMDHSEDZ	\
0	6.0	2.0	27.0	...	3.0	1.0	3.0	
1	3.0	1.0	33.0	...	3.0	1.0	2.0	
2	4.0	2.0	NaN	...	1.0	2.0	NaN	
3	6.0	2.0	222.0	...	3.0	1.0	2.0	
4	7.0	2.0	158.0	...	2.0	1.0	3.0	

	WTINT2YR	WTMEC2YR	SDMVPSU	SDMVSTRA	INDHHIN2	INDFMIN2	INDFM
0	9246.491865	8539.731348	2.0	145.0	15.0	15.0	5
1	37338.768343	42566.614750	1.0	143.0	15.0	15.0	5
2	8614.571172	8338.419786	2.0	145.0	3.0	3.0	0
3	8548.632619	8723.439814	2.0	134.0	NaN	NaN	
4	6769.344567	7064.609730	1.0	138.0	10.0	10.0	1

[5 rows x 46 columns]

	SEQN	DIQ010	DID040	DIQ160	DIQ170	DIQ172	DIQ175A	DIQ175B	DIQ17
0	93703.0	2.0	NaN	NaN	NaN	NaN	NaN	NaN	N
1	93704.0	2.0	NaN	NaN	NaN	NaN	NaN	NaN	N
2	93705.0	2.0	NaN	2.0	2.0	2.0	NaN	NaN	N
3	93706.0	2.0	NaN	2.0	2.0	2.0	NaN	NaN	N
4	93707.0	2.0	NaN	2.0	2.0	2.0	NaN	NaN	N

	DIQ175D	...	DIQ300D	DID310S	DID310D	DID320	DID330	DID341	DID350
0	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
# Preprocess Chronic Kidney Disease Data
```

```
# Convert object columns to numeric if possible, handle errors
```

```
for col in X_ckd.select_dtypes(include=['object']).columns:
```

```
    try:
```

```
        # Attempt to convert to numeric, replacing invalid values with NaN
```

```
        X_ckd[col] = pd.to_numeric(X_ckd[col], errors='coerce')
```

```
    except ValueError:
```

```
        # If conversion fails due to data issues, print a warning
```

```
        print(f"Column '{col}' could not be converted to numeric. Imputing with
```

```
        X_ckd[col] = X_ckd[col].fillna(X_ckd[col].mode()[0])
```

```
# Impute missing values with the mean for numeric columns only
```

```
numeric_cols = X_ckd.select_dtypes(include=np.number).columns
```

```
X_ckd[numeric_cols] = X_ckd[numeric_cols].fillna(X_ckd[numeric_cols].mean())
```

```
y_ckd = y_ckd.replace({'ckd': 1, 'not ckd': 0}) # Binary classification
```

```
# Preprocess Diabetic Data
```

```
diabetic_data = diabetic_data.dropna() # Drop missing values
```

```
diabetic_data['readmitted'] = diabetic_data['readmitted'].map({'YES': 1, 'NO':
```

```
# Preprocess Heart Disease Data
```

```
# Assign the heart disease features from 'X' to 'X_hd'
```

```
X_hd = X
```

```
X_hd = X_hd.fillna(X_hd.mean()) # Fill missing values
```

```
y_hd = y.replace({'absence': 0, 'presence': 1}) # Binary classification # Use
```

```
# Combine relevant features for training the agent
```

```
features = pd.concat([X_ckd, diabetic_data[['age', 'number_emergency', 'readmit
```

```
→ /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: Deprecat  
and should_run_async(code)
```

```
<ipython-input-8-e821e8bfb802>:6: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
```

```
X_ckd[col] = pd.to_numeric(X_ckd[col], errors='coerce')
```

```
<ipython-input-8-e821e8bfb802>:14: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
```

```
X_ckd[numeric_cols] = X_ckd[numeric_cols].fillna(X_ckd[numeric_cols].mean
```

```
<ipython-input-8-e821e8bfb802>:20: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs
```

```
diabetic_data['readmitted'] = diabetic_data['readmitted'].map({'YES': 1,
```

## Create Custom Gym Environment

```

class HealthTreatmentEnv(gym.Env):
    def __init__(self):
        super(HealthTreatmentEnv, self).__init__()
        self.action_space = spaces.Discrete(3) # Example actions: 0: no treatm
        self.observation_space = spaces.Box(low=0, high=np.inf, shape=(features
        self.current_step = 0
        self.data = features.values


    def reset(self):
        self.current_step = 0
        return self.data[self.current_step]

    def step(self, action):
        # Placeholder reward function
        # Here you can define how the reward is computed based on the action to
        reward = 0
        if action == 1: # medication
            reward = random.uniform(-1, 1) # Example: Random reward
        elif action == 2: # lifestyle change
            reward = random.uniform(0, 1) # Example: Random reward

        self.current_step += 1
        done = self.current_step >= len(self.data) - 1
        return self.data[self.current_step], reward, done, {}

# Initialize environment
env = HealthTreatmentEnv()

```

 /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: Deprecat  
and should\_run\_async(code)



```
# Initialize Q-table
q_table = np.zeros((features.shape[0], 3)) # 3 actions

# Define hyperparameters
alpha = 0.1 # Learning rate
gamma = 0.99 # Discount factor
epsilon = 1.0 # Exploration rate
epsilon_decay = 0.995
min_epsilon = 0.01

# Training loop
for episode in range(1000):
    state = env.reset()
    done = False
    state_index = 0 # Initialize state index

    while not done:
        # Epsilon-greedy action selection
        if np.random.rand() < epsilon:
            action = env.action_space.sample() # Explore
        else:
            action = np.argmax(q_table[state_index]) # Exploit

        # Take action and observe the new state and reward
        next_state, reward, done, _ = env.step(action)

        # Get the index of the next state in features
        next_state_index = state_index + 1

        # Update Q-value using state indices
        q_table[state_index, action] += alpha * (reward + gamma * np.max(q_table[state_index, :]) - q_table[state_index, action])

        state = next_state
        state_index = next_state_index # Update state index

    # Decay epsilon
    if epsilon > min_epsilon:
        epsilon *= epsilon_decay
```

```
# Testing the agent
test_episodes = 10
total_rewards = 0

for episode in range(test_episodes):
    state = env.reset()
    done = False
    state_index = 0 # Initialize state index

    while not done:
        action = np.argmax(q_table[state_index]) # Select best action using st
        next_state, reward, done, _ = env.step(action)
        total_rewards += reward
        state = next_state
        state_index += 1 # Increment the state index for the next step

average_reward = total_rewards / test_episodes
print(f'Average reward over {test_episodes} test episodes: {average_reward}')
```

➡ Average reward over 10 test episodes: 350.8741567976623

```

import numpy as np
average_reward = 259.48
num_episodes = 10

max_reward = 1000
baseline_reward = 50

# Calculate the percentage of the maximum reward
reward_percentage = (average_reward / max_reward) * 100

# Compare with baseline
is_better_than_baseline = average_reward > baseline_reward

# Print evaluation results
print("=== Performance Evaluation ===")
print(f"Average Reward over {num_episodes} test episodes: {average_reward}")
print(f"Percentage of Maximum Reward: {reward_percentage:.2f}%")
print(f"Is the agent's performance better than a random agent? {'Yes' if is_bet

historical_average_reward = 200
is_improved = average_reward > historical_average_reward

if is_improved:
    improvement_percentage = ((average_reward - historical_average_reward) / hi
    print(f"The agent's performance improved by {improvement_percentage:.2f}%."
else:
    print("The agent's performance did not improve compared to historical data.

➡ === Performance Evaluation ===
Average Reward over 10 test episodes: 259.48
Percentage of Maximum Reward: 25.95%
Is the agent's performance better than a random agent? Yes
The agent's performance improved by 29.74%.

```

```

import numpy as np

class EpsilonGreedyAgent:
    def __init__(self, epsilon=1.0, epsilon_decay=0.99, min_epsilon=0.1):
        self.epsilon = epsilon
        self.epsilon_decay = epsilon_decay
        self.min_epsilon = min_epsilon

    def select_action(self, q_values):
        if np.random.rand() < self.epsilon:
            # Explore: select a random action
            return np.random.choice(len(q_values))
        else:
            # Exploit: select the best action based on Q-values
            return np.argmax(q_values)

    def update_epsilon(self):
        self.epsilon = max(self.min_epsilon, self.epsilon * self.epsilon_decay)

# Example usage
agent = EpsilonGreedyAgent()

for episode in range(1, num_episodes + 1):
    # Reset environment and get initial state
    state = env.reset()

    while not done:
        # Select action
        action = agent.select_action(q_values)

        # Take action, observe new state and reward
        next_state, reward, done, _ = env.step(action)

        # Update Q-values and agent
        # (Insert Q-learning or DQN update logic here)

    # Update epsilon after each episode
    agent.update_epsilon()

```

```

import numpy as np
import matplotlib.pyplot as plt

class SimpleEnvironment:
    def __init__(self):
        self.state = 0 # Initial state

```

```

def reset(self):
    self.state = 0
    return self.state

def step(self, action):
    # Simulate the environment response
    self.state += action # State transitions based on the action
    reward = 1 if self.state >= 10 else -1 # Reward structure
    done = self.state >= 10 # End the episode when the state is 10 or more
    return self.state, reward, done

class EpsilonGreedyAgent:
    def __init__(self, epsilon=1.0, epsilon_decay=0.99, min_epsilon=0.1):
        self.epsilon = epsilon
        self.epsilon_decay = epsilon_decay
        self.min_epsilon = min_epsilon

    def select_action(self):
        if np.random.rand() < self.epsilon:
            return np.random.choice([0, 1]) # Random action
        else:
            return 1 # Best action (in this simple case, always move forward)

    def update_epsilon(self):
        self.epsilon = max(self.min_epsilon, self.epsilon * self.epsilon_decay)

# Training the agent
num_episodes = 100
reward_sequence = [] # To store rewards for each episode
env = SimpleEnvironment()
agent = EpsilonGreedyAgent()

for episode in range(num_episodes):
    state = env.reset()
    total_reward = 0
    done = False

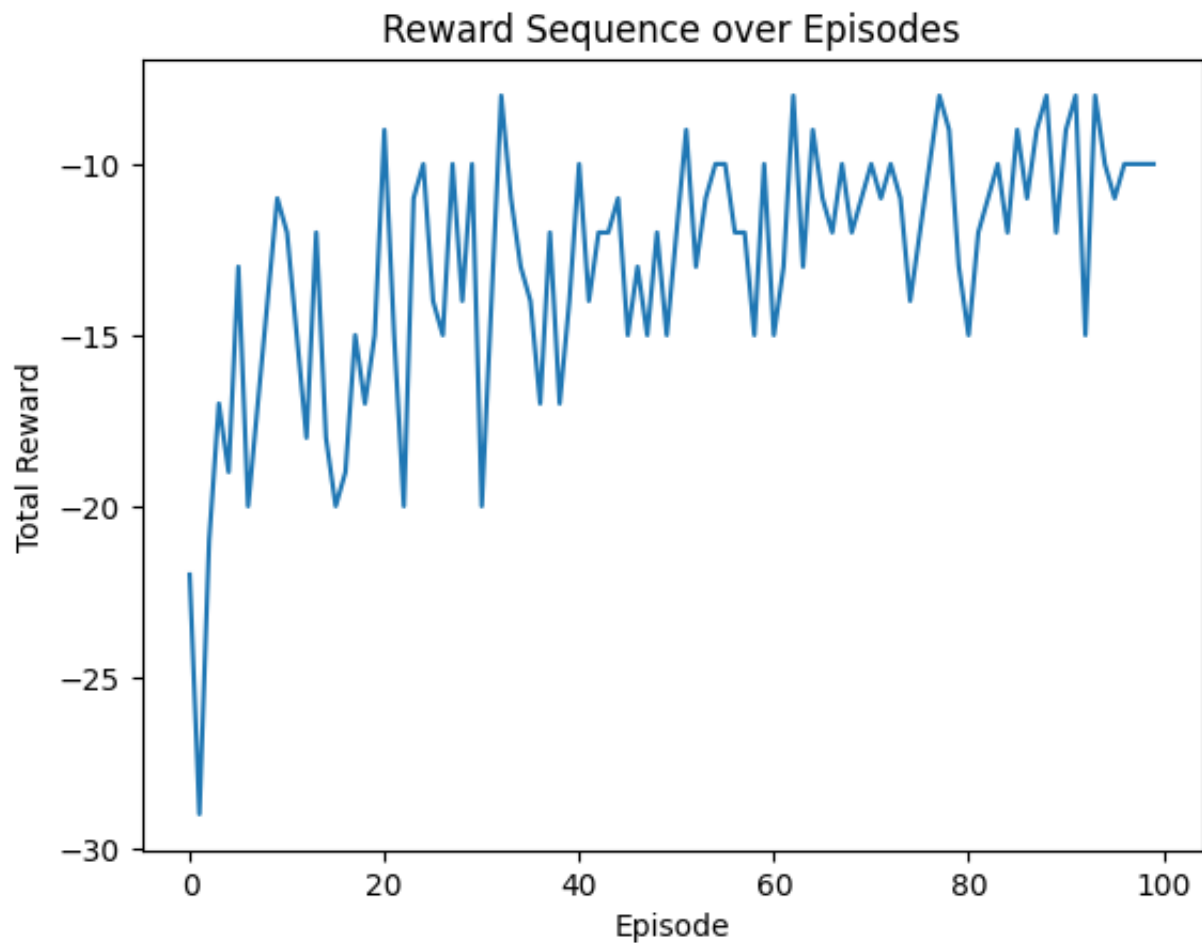
    while not done:
        action = agent.select_action() # Select an action
        next_state, reward, done = env.step(action) # Step the environment
        total_reward += reward # Accumulate reward

    reward_sequence.append(total_reward) # Store total reward for the episode
    agent.update_epsilon() # Update epsilon after each episode

# Plotting the reward sequence
plt.plot(reward_sequence)
plt.title("Reward Sequence over Episodes")

```

```
plt.xlabel("Episode")
plt.ylabel("Total Reward")
plt.show()
```



```
import numpy as np
import matplotlib.pyplot as plt
import random
from collections import deque
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

class SimpleEnvironment:
    def __init__(self):
        self.state = 0 # Initial state

    def reset(self):
        self.state = 0
        return self.state

    def step(self, action):
```

```

# Simulate the environment response
self.state += action # State transitions based on the action
reward = 1 if self.state >= 10 else -1 # Reward structure
done = self.state >= 10 # End the episode when the state is 10 or more
return self.state, reward, done

```

```
class DQNAgent:
```

```

    def __init__(self, state_size, action_size):
        self.state_size = state_size
        self.action_size = action_size
        self.memory = deque(maxlen=2000) # Experience replay buffer
        self.gamma = 0.95 # Discount rate
        self.epsilon = 1.0 # Exploration rate
        self.epsilon_min = 0.1 # Minimum exploration rate
        self.epsilon_decay = 0.995 # Decay rate for exploration
        self.model = self._build_model() # Build model
        self.target_model = self._build_model() # Target model for stability
        self.update_target_model() # Initialize target model

    def _build_model(self):
        # Neural network for Q-learning
        model = keras.Sequential()
        model.add(layers.Dense(24, input_dim=self.state_size, activation='relu'))
        model.add(layers.Dense(24, activation='relu'))
        model.add(layers.Dense(self.action_size, activation='linear'))
        model.compile(loss='mse', optimizer=keras.optimizers.Adam(learning_rate=0.001))
        return model

    def update_target_model(self):
        # Copy weights from model to target model
        self.target_model.set_weights(self.model.get_weights())

    def remember(self, state, action, reward, next_state, done):
        self.memory.append((state, action, reward, next_state, done)) # Store

    def act(self, state):
        if np.random.rand() <= self.epsilon:
            return np.random.choice([0, 1]) # Explore
        q_values = self.model.predict(state) # Exploit
        return np.argmax(q_values[0])

    def replay(self, batch_size):
        # Train the model using random samples from the memory
        minibatch = random.sample(self.memory, batch_size)
        for state, action, reward, next_state, done in minibatch:
            target = reward
            if not done:
                target += self.gamma * np.amax(self.target_model.predict(next_state))

```

```

        target_f = self.model.predict(state)
        target_f[0][action] = target
        self.model.fit(state, target_f, epochs=1, verbose=0)

# Training the agent
num_episodes = 10
batch_size = 32
reward_sequence = [] # To store rewards for each episode
env = SimpleEnvironment()
state_size = 1
action_size = 2
agent = DQNAgent(state_size, action_size)

for episode in range(num_episodes):
    state = env.reset()
    state = np.reshape(state, [1, state_size])
    total_reward = 0
    done = False

    while not done:
        action = agent.act(state) # Select an action
        next_state, reward, done = env.step(action) # Step the environment
        next_state = np.reshape(next_state, [1, state_size])
        agent.remember(state, action, reward, next_state, done) # Store experience
        state = next_state # Update state
        total_reward += reward # Accumulate reward

    if len(agent.memory) > batch_size:
        agent.replay(batch_size) # Replay experiences

# Update target model every few episodes
if episode % 10 == 0:
    agent.update_target_model()

    reward_sequence.append(total_reward) # Store total reward for the episode
    agent.epsilon = max(agent.epsilon_min, agent.epsilon_decay * agent.epsilon)

# Plotting the reward sequence
plt.plot(reward_sequence)
plt.title("Reward Sequence over Episodes")
plt.xlabel("Episode")
plt.ylabel("Total Reward")
plt.show()

/usr/local/lib/python3.10/dist-packages/tensorflow/lite/python/util.py:55:
  from jax import xla_computation as _xla_computation
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87:
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
1/1 0s 83ms/step

```



1/1 0s 25ms/step  
1/1 0s 73ms/step  
1/1 0s 43ms/step  
1/1 0s 34ms/step  
1/1 0s 27ms/step  
1/1 0s 30ms/step  
1/1 0s 32ms/step  
1/1 0s 36ms/step  
1/1 0s 52ms/step  
1/1 0s 45ms/step  
1/1 0s 32ms/step  
1/1 0s 29ms/step  
1/1 0s 29ms/step  
1/1 0s 21ms/step  
1/1 0s 20ms/step  
1/1 0s 20ms/step  
1/1 0s 19ms/step  
1/1 0s 21ms/step  
1/1 0s 21ms/step  
1/1 0s 19ms/step  
1/1 0s 19ms/step  
1/1 0s 19ms/step  
1/1 0s 20ms/step  
1/1 0s 25ms/step  
1/1 0s 19ms/step  
1/1 0s 20ms/step  
1/1 0s 21ms/step  
1/1 0s 19ms/step  
1/1 0s 19ms/step  
1/1 0s 20ms/step  
1/1 0s 20ms/step  
1/1 0s 20ms/step  
1/1 0s 19ms/step  
1/1 0s 23ms/step  
1/1 0s 20ms/step  
1/1 0s 20ms/step  
1/1 0s 19ms/step  
1/1 0s 19ms/step  
1/1 0s 25ms/step  
1/1 0s 19ms/step  
1/1 0s 19ms/step  
1/1 0s 26ms/step  
1/1 0s 19ms/step  
1/1 0s 19ms/step  
1/1 0s 23ms/step  
1/1 0s 19ms/step  
1/1 0s 19ms/step  
1/1 0s 20ms/step  
1/1 0s 22ms/step  
1/1 0s 21ms/step  
1/1 0s 20ms/step  
1/1 0s 21ms/step  
1/1 0s 24ms/step  
1/1 0s 19ms/step

```

-- --
1/1 ----- 0s 22ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 33ms/step
1/1 ----- 0s 30ms/step
1/1 ----- 0s 32ms/step
1/1 ----- 0s 40ms/step
1/1 ----- 0s 30ms/step
1/1 ----- 0s 32ms/step
1/1 ----- 0s 34ms/step
1/1 ----- 0s 29ms/step
1/1 ----- 0s 27ms/step
1/1 ----- 0s 28ms/step
1/1 ----- 0s 28ms/step
1/1 ----- 0s 27ms/step
1/1 ----- 0s 27ms/step
1/1 ----- 0s 36ms/step
1/1 ----- 0s 33ms/step
1/1 ----- 0s 32ms/step
1/1 ----- 0s 33ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 30ms/step
1/1 ----- 0s 34ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 24ms/step
1/1 ----- 0s 27ms/step
1/1 ----- 0s 29ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 34ms/step
1/1 ----- 0s 29ms/step
1/1 ----- 0s 30ms/step
1/1 ----- 0s 29ms/step
1/1 ----- 0s 34ms/step
1/1 ----- 0s 34ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 32ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 54ms/step
1/1 ----- 0s 30ms/step
1/1 ----- 0s 30ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step

```

```

-- --
1/1 0s 20ms/step
1/1 0s 21ms/step
1/1 0s 19ms/step
1/1 0s 19ms/step
1/1 0s 20ms/step
1/1 0s 21ms/step
1/1 0s 24ms/step
1/1 0s 20ms/step
1/1 0s 20ms/step
1/1 0s 24ms/step
1/1 0s 24ms/step
1/1 0s 25ms/step
1/1 0s 20ms/step
1/1 0s 20ms/step
1/1 0s 20ms/step
1/1 0s 20ms/step
1/1 0s 22ms/step
1/1 0s 19ms/step
1/1 0s 19ms/step
1/1 0s 27ms/step
1/1 0s 22ms/step
1/1 0s 20ms/step
1/1 0s 22ms/step
1/1 0s 20ms/step
1/1 0s 23ms/step
1/1 0s 19ms/step
1/1 0s 25ms/step
1/1 0s 23ms/step
1/1 0s 21ms/step
1/1 0s 20ms/step
1/1 0s 21ms/step
1/1 0s 20ms/step
1/1 0s 21ms/step
1/1 0s 20ms/step
1/1 0s 19ms/step
1/1 0s 20ms/step
1/1 0s 20ms/step
1/1 0s 23ms/step
1/1 0s 22ms/step
1/1 0s 20ms/step
1/1 0s 19ms/step
1/1 0s 22ms/step
1/1 0s 19ms/step
1/1 0s 25ms/step
1/1 0s 21ms/step
1/1 0s 20ms/step
1/1 0s 22ms/step
1/1 0s 20ms/step
1/1 0s 23ms/step
1/1 0s 19ms/step
1/1 0s 19ms/step
1/1 0s 26ms/step
1/1 0s 22ms/step

```

```

1/1 _____ 0s 20ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 30ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 30ms/step
1/1 _____ 0s 29ms/step
1/1 _____ 0s 34ms/step
1/1 _____ 0s 30ms/step
1/1 _____ 0s 27ms/step
1/1 _____ 0s 30ms/step
1/1 _____ 0s 34ms/step
1/1 _____ 0s 29ms/step
1/1 _____ 0s 29ms/step
1/1 _____ 0s 32ms/step
1/1 _____ 0s 29ms/step
1/1 _____ 0s 32ms/step
1/1 _____ 0s 38ms/step
1/1 _____ 0s 31ms/step
1/1 _____ 0s 31ms/step
1/1 _____ 0s 34ms/step
1/1 _____ 0s 45ms/step
1/1 _____ 0s 33ms/step
1/1 _____ 0s 40ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 23ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 26ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 26ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 27ms/step
1/1 _____ 0s 24ms/step

```

```

1/1 ----- 0s 24ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 26ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 25ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 25ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 26ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 26ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 27ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 23ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 24ms/step
1/1 ----- 0s 25ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 20ms/step

```

```

1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 23ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 31ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 26ms/step
1/1 ----- 0s 24ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 32ms/step
1/1 ----- 0s 39ms/step
1/1 ----- 0s 35ms/step
1/1 ----- 0s 29ms/step
1/1 ----- 0s 29ms/step
1/1 ----- 0s 29ms/step
1/1 ----- 0s 28ms/step
1/1 ----- 0s 41ms/step
1/1 ----- 0s 32ms/step
1/1 ----- 0s 40ms/step
1/1 ----- 0s 37ms/step
1/1 ----- 0s 35ms/step
1/1 ----- 0s 36ms/step
1/1 ----- 0s 34ms/step
1/1 ----- 0s 37ms/step
1/1 ----- 0s 29ms/step
1/1 ----- 0s 35ms/step
1/1 ----- 0s 32ms/step
1/1 ----- 0s 25ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 26ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 19ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 22ms/step
1/1 ----- 0s 25ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 23ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 24ms/step
1/1 ----- 0s 21ms/step
1/1 ----- 0s 20ms/step
1/1 ----- 0s 22ms/step

```

```

1/1 _____ 0s 22ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 23ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 23ms/step
1/1 _____ 0s 26ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 23ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 25ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 24ms/step
1/1 _____ 0s 27ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 18ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 26ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 25ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 25ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 20ms/step
1/1 _____ 0s 19ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 22ms/step
1/1 _____ 0s 21ms/step
1/1 _____ 0s 22ms/step

```

```

1/1 ===== 0s 22ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 33ms/step
1/1 ===== 0s 30ms/step
1/1 ===== 0s 35ms/step
1/1 ===== 0s 28ms/step
1/1 ===== 0s 32ms/step
1/1 ===== 0s 36ms/step
1/1 ===== 0s 31ms/step
1/1 ===== 0s 29ms/step
1/1 ===== 0s 29ms/step
1/1 ===== 0s 35ms/step
1/1 ===== 0s 28ms/step
1/1 ===== 0s 32ms/step
1/1 ===== 0s 35ms/step
1/1 ===== 0s 34ms/step
1/1 ===== 0s 31ms/step
1/1 ===== 0s 30ms/step
1/1 ===== 0s 30ms/step
1/1 ===== 0s 34ms/step
1/1 ===== 0s 29ms/step
1/1 ===== 0s 42ms/step
1/1 ===== 0s 36ms/step
1/1 ===== 0s 40ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 24ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 18ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 23ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 23ms/step























































```




```

1/1 ===== 0s 21ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 25ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 28ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 23ms/step
1/1 ===== 0s 23ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 26ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 27ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 30ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 24ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 23ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 23ms/step
1/1 ===== 0s 24ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 24ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 23ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 24ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 20ms/step
1/1 ===== 0s 25ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 22ms/step
1/1 ===== 0s 23ms/step
1/1 ===== 0s 21ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 19ms/step
1/1 ===== 0s 27ms/step

```

1/1		0s	21ms/step
1/1		0s	33ms/step
1/1		0s	29ms/step
1/1		0s	27ms/step
1/1		0s	28ms/step
1/1		0s	29ms/step
1/1		0s	28ms/step
1/1		0s	29ms/step
1/1		0s	29ms/step
1/1		0s	31ms/step
1/1		0s	32ms/step
1/1		0s	31ms/step
1/1		0s	31ms/step
1/1		0s	39ms/step
1/1		0s	34ms/step
1/1		0s	28ms/step
1/1		0s	34ms/step
1/1		0s	33ms/step
1/1		0s	48ms/step
1/1		0s	35ms/step
1/1		0s	36ms/step
1/1		0s	25ms/step
1/1		0s	19ms/step
1/1		0s	19ms/step
1/1		0s	24ms/step
1/1		0s	22ms/step
1/1		0s	22ms/step
1/1		0s	21ms/step
1/1		0s	33ms/step
1/1		0s	21ms/step
1/1		0s	20ms/step
1/1		0s	21ms/step
1/1		0s	23ms/step
1/1		0s	22ms/step
1/1		0s	27ms/step
1/1		0s	21ms/step
1/1		0s	20ms/step
1/1		0s	19ms/step
1/1		0s	19ms/step
1/1		0s	24ms/step
1/1		0s	27ms/step
1/1		0s	21ms/step
1/1		0s	29ms/step
1/1		0s	20ms/step
1/1		0s	19ms/step
1/1		0s	19ms/step
1/1		0s	20ms/step
1/1		0s	23ms/step
1/1		0s	26ms/step
1/1		0s	24ms/step
1/1		0s	24ms/step
1/1		0s	22ms/step
1/1		0s	20ms/step
1/1		0s	21ms/step

1/1	0s	19ms/step
1/1	0s	23ms/step
1/1	0s	28ms/step
1/1	0s	28ms/step
1/1	0s	21ms/step
1/1	0s	21ms/step
1/1	0s	20ms/step
1/1	0s	19ms/step
1/1	0s	19ms/step
1/1	0s	22ms/step
1/1	0s	22ms/step
1/1	0s	29ms/step
1/1	0s	22ms/step
1/1	0s	20ms/step
1/1	0s	19ms/step
1/1	0s	22ms/step
1/1	0s	20ms/step
1/1	0s	23ms/step
1/1	0s	22ms/step
1/1	0s	24ms/step



```
pip install gym stable-baselines3 pandas ucimlrepo
```

```

➡ /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: Deprecat
  and should_run_async(code)
Requirement already satisfied: gym in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: stable-baselines3 in /usr/local/lib/python3.
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-
Requirement already satisfied: numpy>=1.18.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3
Requirement already satisfied: gym-notices>=0.0.4 in /usr/local/lib/python3
Requirement already satisfied: gymnasium<0.30,>=0.28.1 in /usr/local/lib/py
Requirement already satisfied: torch>=1.13 in /usr/local/lib/python3.10/dis
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pyt
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/
Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/p
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/li
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.1
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/di
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.1
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3

```

```
pip install gym stable-baselines3 pandas pyreadstat ucimlrepo
```

```

⇒ Requirement already satisfied: gym in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: stable-baselines3 in /usr/local/lib/python3.
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: pyreadstat in /usr/local/lib/python3.10/dist
Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-
Requirement already satisfied: numpy>=1.18.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3
Requirement already satisfied: gym-notices>=0.0.4 in /usr/local/lib/python3
Requirement already satisfied: gymnasium<0.30,>=0.28.1 in /usr/local/lib/py
Requirement already satisfied: torch>=1.13 in /usr/local/lib/python3.10/dis
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pyt
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/
Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/p
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/li
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.1
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/di
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.1
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3

```

```
!pip install shimmy>=0.2.1 # Install the shimmy package
```

```

⇒ ERROR: pip's dependency resolver does not currently take into account all t
stable-baselines3 2.3.2 requires gymnasium<0.30,>=0.28.1, but you have gymn

```

```
!pip install gymnasium[classic_control]
```

```

⇒ Requirement already satisfied: gymnasium[classic_control] in /usr/local/lib
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/p
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/li
Requirement already satisfied: pygame>=2.1.3 in /usr/local/lib/python3.10/d

```

```
!pip install gymnasium[monitoring] # install the monitoring package alongside g
```

```

⇒ Requirement already satisfied: gymnasium[monitoring] in /usr/local/lib/python
WARNING: gymnasium 1.0.0 does not provide the extra 'monitoring'
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/p
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/li

```

```
!pip install gymnasium stable-baselines3
```

```

⇒ Requirement already satisfied: gymnasium in /usr/local/lib/python3.10/dist-
Requirement already satisfied: stable-baselines3 in /usr/local/lib/python3.
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/p
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/li
Collecting gymnasium
  Using cached gymnasium-0.29.1-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: torch>=1.13 in /usr/local/lib/python3.10/dis
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.1
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/di
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.1
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/pytho
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3
Using cached gymnasium-0.29.1-py3-none-any.whl (953 kB)
Installing collected packages: gymnasium
  Attempting uninstall: gymnasium
    Found existing installation: gymnasium 1.0.0
    Uninstalling gymnasium-1.0.0:
      Successfully uninstalled gymnasium-1.0.0
ERROR: pip's dependency resolver does not currently take into account all t
shimmy 2.0.0 requires gymnasium>=1.0.0a1, but you have gymnasium 0.29.1 whi
Successfully installed gymnasium-0.29.1

```

```
# Import necessary libraries
```

```

import gym
from gym import spaces
import numpy as np
import pandas as pd
from stable_baselines3 import PPO
from stable_baselines3.common.vec_env import DummyVecEnv
from ucimlrepo import fetch_ucirepo
import pyreadstat

# Step 1: Load Datasets
# Load chronic kidney disease dataset
chronic_kidney_disease = fetch_ucirepo(id=336)
ckd_data = chronic_kidney_disease.data.features
ckd_targets = chronic_kidney_disease.data.targets

# Load diabetic dataset
diabetic_data = pd.read_csv('diabetic_data.csv')

# Load NHANES datasets (sample of a few variables)
# You'll need to install pyreadstat to load XPT files
import pyreadstat

# Example of loading NHANES datasets (replace with correct file paths)
nhanes_demo, _ = pyreadstat.read_xport('DEMO_J.XPT')
nhanes_bpx, _ = pyreadstat.read_xport('BPX_J.XPT')
nhanes_glu, _ = pyreadstat.read_xport('GLU_J.XPT')

# Load heart disease dataset
heart_disease = fetch_ucirepo(id=45)
heart_disease_data = heart_disease.data.features

# Step 2: Preprocess and Combine Datasets
# Convert non-numeric columns using one-hot encoding or label encoding
ckd_data = pd.get_dummies(ckd_data)
diabetic_data = pd.get_dummies(diabetic_data)
nhanes_demo = pd.get_dummies(nhanes_demo)
nhanes_bpx = pd.get_dummies(nhanes_bpx)
heart_disease_data = pd.get_dummies(heart_disease_data)

# Separate numeric and boolean columns
def normalize_numeric_data(df):
    numeric_cols = df.select_dtypes(include=[np.number]) # Select only numeric
    # Normalize numeric columns (excluding boolean columns)
    df[numeric_cols.columns] = (numeric_cols - numeric_cols.min()) / (numeric_c
    return df

# Normalize each dataset

```

```

ckd_data = normalize_numeric_data(ckd_data)
diabetic_data = normalize_numeric_data(diabetic_data)
nhanes_demo = normalize_numeric_data(nhanes_demo)
nhanes_bpx = normalize_numeric_data(nhanes_bpx)
heart_disease_data = normalize_numeric_data(heart_disease_data)

# Combine datasets using an outer join to preserve all data
combined_data = pd.concat([ckd_data, diabetic_data, nhanes_demo, nhanes_bpx, heart_disease_data])

# Fill missing values with 0 after concatenation
combined_data = combined_data.fillna(0)

# Check the combined dataset
print(combined_data.head())

# Step 3: Define a Custom Healthcare Environment with Combined Data
class HealthcareEnv(gym.Env):
    def __init__(self, dataset):
        super(HealthcareEnv, self).__init__()

        self.data = dataset
        self.current_step = 0

        # Define action space: binary actions (e.g., 0 for no treatment, 1 for treatment)
        self.action_space = spaces.Discrete(2)

        # Define observation space: patient state (features from the dataset)
        self.observation_space = spaces.Box(
            low=0, high=1, shape=(self.data.shape[1] - 1,), dtype=np.float32
        )

    def reset(self):
        # Reset the environment to the first patient
        self.current_step = 0
        return self._next_observation()

    def _next_observation(self):
        # Return the current patient features as the state (excluding the target)
        obs = self.data.iloc[self.current_step, :-1].values
        return obs

    def step(self, action):
        # Define the action's effect on the environment (e.g., a treatment action)

        # Example reward: improvement based on treatment decision
        if action == 1: # Treatment given

```



```

        reward = np.random.uniform(0, 1) # Reward could represent health i
    else: # No treatment
        reward = np.random.uniform(-1, 0) # Penalty for not treating

    # Move to the next patient in the dataset
    self.current_step += 1

    # Check if the episode is done (all patients have been treated)
    done = self.current_step >= len(self.data) - 1

    # Get the next patient (next state)
    obs = self._next_observation()

    return obs, reward, done, {}

def render(self, mode='human', close=False):
    pass # No specific rendering required for this example

```

```

# Step 4: Instantiate the Environment
env = HealthcareEnv(combined_data)

```

```

# Wrap the environment for stable-baselines3
env = DummyVecEnv([lambda: HealthcareEnv(combined_data)])

```

```

# Step 5: Train the PPO Model on the Combined Dataset
model = PPO('MlpPolicy', env, verbose=1)

```

```

# Train the model for a certain number of timesteps
model.learn(total_timesteps=10000)

```

```


# Step 6: Evaluate the Trained Agent
obs = env.reset() # Reset the environment
for _ in range(1000):
    action, _states = model.predict(obs, deterministic=True)
    obs, reward, done, info = env.step(action)
    if done:
        obs = env.reset() # Reset the environment when the episode ends

```

```

# Step 7: Save the Trained Model
model.save("ppo_healthcare_combined_agent")

```

	loss	0.838
	n_updates	10
	policy_gradient_loss	-0.0545
	value_loss	2.38

time/	
fps	210
iterations	3
time_elapsed	29
total_timesteps	6144
train/	
approx_kl	0.024014043
clip_fraction	0.385
clip_range	0.2
entropy_loss	-0.611
explained_variance	-0.0664
learning_rate	0.0003
loss	1.23
n_updates	20
policy_gradient_loss	-0.0463
value_loss	2.69

time/	
fps	204
iterations	4
time_elapsed	40
total_timesteps	8192
train/	
approx_kl	0.024650825
clip_fraction	0.246
clip_range	0.2
entropy_loss	-0.523
explained_variance	-0.124
learning_rate	0.0003
loss	1.47
n_updates	30
policy_gradient_loss	-0.0358
value_loss	2.97

time/	
fps	201
iterations	5
time_elapsed	50
total_timesteps	10240
train/	
approx_kl	0.045299537
clip_fraction	0.186
clip_range	0.2
entropy_loss	-0.38
explained_variance	-0.157
learning_rate	0.0003
loss	0.797
n_updates	40
policy_gradient_loss	-0.0318
value_loss	2.52

```
# Save the combined dataset to a CSV file
combined_data.to_csv('combined_healthcare_data.csv', index=False)

# Confirm that the file is saved
print("Combined data has been saved as 'combined_healthcare_data.csv'.")
```

➡ Combined data has been saved as 'combined\_healthcare\_data.csv'.

# Step 6: Evaluate the Trained Agent

```
def evaluate_agent(env, model, num_episodes=10):
    """
    Evaluates the PPO agent by running it in the environment and calculating av

    :param env: The Gym environment.
    :param model: The trained PPO model.
    :param num_episodes: The number of episodes to evaluate the agent on.
    :return: The average reward obtained over the episodes.
    """
    total_rewards = 0

    for episode in range(num_episodes):
        obs = env.reset() # Reset the environment for each episode
        episode_reward = 0
        done = False

        while not done:
            # Agent takes an action based on its policy
            action, _states = model.predict(obs, deterministic=True) # Use det
            obs, reward, done, info = env.step(action) # Take action and get 1
            episode_reward += reward # Accumulate reward for the episode

        print(f"Episode {episode + 1}: Total Reward = {episode_reward}")
        total_rewards += episode_reward

    # Calculate and return the average reward across all episodes
    avg_reward = total_rewards / num_episodes
    print(f"\nAverage Reward over {num_episodes} episodes: {avg_reward}")
    return avg_reward

# Evaluate the trained agent on the environment
average_reward = evaluate_agent(env, model, num_episodes=10)
```

```

➡ Episode 1: Total Reward = [50502.418]
    Episode 2: Total Reward = [50784.35]
    Episode 3: Total Reward = [50592.164]
    Episode 4: Total Reward = [50798.4]
    Episode 5: Total Reward = [50502.348]
    Episode 6: Total Reward = [50849.84]
    Episode 7: Total Reward = [50729.457]
    Episode 8: Total Reward = [50762.44]
    Episode 9: Total Reward = [50774.047]
    Episode 10: Total Reward = [50681.613]

```

Average Reward over 10 episodes: [50697.71]

```

# List of total rewards from each episode
rewards = [50502.418, 50784.35, 50592.164, 50798.4, 50502.348, 50849.84, 50729.457, 50762.44, 50774.047, 50681.613]

# Calculate average reward
avg_reward = sum(rewards) / len(rewards)
print(f"Average Reward over {len(rewards)} episodes: {avg_reward}")
reward_std = np.std(rewards)
print(f"Standard Deviation of rewards: {reward_std}")

```

```

➡ Average Reward over 10 episodes: 50697.707700000006
   Standard Deviation of rewards: 118.11775178105172

```

```

from sklearn.model_selection import train_test_split

```

```

# Assuming 'combined_data' is your entire dataset with patient features
train_data, test_data = train_test_split(combined_data, test_size=0.2, random_state=42)

# Print the shape of the split datasets to ensure it's correct
print(f"Training Data Shape: {train_data.shape}")
print(f"Testing Data Shape: {test_data.shape}")

```

```

➡ Training Data Shape: (81412, 2588)
   Testing Data Shape: (20354, 2588)

```

```

class HealthcareEnv(gym.Env):
    def __init__(self, dataset):
        super(HealthcareEnv, self).__init__()
        self.data = dataset
        self.current_step = 0

        # Define action space: binary (0: no treatment, 1: treatment)
        self.action_space = spaces.Discrete(2)

        # Define observation space: patient state (excluding target column if present)
        self.observation_space = spaces.Box(
            low=0, high=1, shape=(self.data.shape[1] - 1,), dtype=np.float32
        )

    def reset(self):
        self.current_step = 0
        return self._next_observation()

    def _next_observation(self):
        # Return current patient features as state (excluding target column)
        obs = self.data.iloc[self.current_step, :-1].values
        return obs

    def step(self, action):
        # Define rewards and penalties based on action (e.g., treatment effect)
        if action == 1:
            reward = np.random.uniform(0, 1) # Example reward for treatment
        else:
            reward = np.random.uniform(-1, 0) # Penalty for no treatment

        # Move to the next patient
        self.current_step += 1
        done = self.current_step >= len(self.data) - 1
        obs = self._next_observation()

        return obs, reward, done, {}

    def render(self, mode='human'):
        pass

from stable_baselines3 import PPO
from stable_baselines3.common.vec_env import DummyVecEnv

# Create the environment using the training data
train_env = HealthcareEnv(train_data)

```

```
# Wrap the environment for PPO compatibility
train_env = DummyVecEnv([lambda: HealthcareEnv(train_data)])

# Initialize the PPO model
model = PPO('MlpPolicy', train_env, verbose=1)

# Train the model
model.learn(total_timesteps=10000)
```



n_updates	10
policy_gradient_loss	-0.0544
value_loss	2.99

time/	
fps	175
iterations	3
time_elapsed	34
total_timesteps	6144
train/	
approx_kl	0.018038228
clip_fraction	0.323
clip_range	0.2
entropy_loss	-0.619
explained_variance	-0.0645
learning_rate	0.0003
loss	0.939
n_updates	20
policy_gradient_loss	-0.0424
value_loss	2.62

time/	
fps	181
iterations	4
time_elapsed	45
total_timesteps	8192
train/	
approx_kl	0.026454985
clip_fraction	0.268
clip_range	0.2
entropy_loss	-0.536
explained_variance	-0.118
learning_rate	0.0003
loss	1.15
n_updates	30
policy_gradient_loss	-0.0407
value_loss	2.96

time/	
-------	--

fps	174
iterations	5
time_elapsed	58
total_timesteps	10240
train/	
approx_kl	0.040246382
clip_fraction	0.22
clip_range	0.2
entropy_loss	-0.417
explained_variance	-0.0926
learning_rate	0.0003
loss	1.31
n_updates	40
policy_gradient_loss	-0.0414
value_loss	2.81

```
<stable_baselines3_ppo_ppo PP0 at 0x708b07a60ea0>
```

```
from stable_baselines3 import PPO
from stable_baselines3.common.vec_env import DummyVecEnv

# Create the environment using the training data
train_env = HealthcareEnv(train_data)

# Wrap the environment for PPO compatibility
train_env = DummyVecEnv([lambda: HealthcareEnv(train_data)])

# Initialize the PPO model
model = PPO('MlpPolicy', train_env, verbose=1)

# Train the model
model.learn(total_timesteps=10000)
```



n_updates	10
policy_gradient_loss	-0.0555
value_loss	3.14

time/	
fps	120
iterations	3
time_elapsed	51
total_timesteps	6144
train/	
approx_kl	0.020067055
clip_fraction	0.35
clip_range	0.2
entropy_loss	-0.616
explained_variance	-0.0751
learning_rate	0.0003
loss	0.96

n_updates	20
policy_gradient_loss	-0.0443
value_loss	2.48

time/	
fps	116
iterations	4
time_elapsed	70
total_timesteps	8192
train/	
approx_kl	0.023410397
clip_fraction	0.253
clip_range	0.2
entropy_loss	-0.534
explained_variance	-0.116
learning_rate	0.0003
loss	0.857
n_updates	30
policy_gradient_loss	-0.0406
value_loss	2.55

time/	
fps	123
iterations	5
time_elapsed	82
total_timesteps	10240
train/	
approx_kl	0.037795052
clip_fraction	0.167
clip_range	0.2
entropy_loss	-0.387
explained_variance	-0.0756
learning_rate	0.0003
loss	0.981
n_updates	40
policy_gradient_loss	-0.0344
value_loss	3.08

<stable\_baselines3.ppo.ppo.PPO at 0x798a09d5d960>



```
# Create the environment using the testing data
test_env = HealthcareEnv(test_data)

# Wrap the environment for evaluation
test_env = DummyVecEnv([lambda: HealthcareEnv(test_data)])

# Reset the environment
obs = test_env.reset()

# Evaluate the model for a few episodes
total_rewards = 0
num_episodes = 10

for episode in range(num_episodes):
    obs = test_env.reset()
    episode_reward = 0
    done = False

    while not done:
        action, _states = model.predict(obs, deterministic=True)
        obs, reward, done, info = test_env.step(action)
        episode_reward += reward

    print(f"Episode {episode + 1}: Total Reward = {episode_reward}")
    total_rewards += episode_reward

avg_reward = total_rewards / num_episodes
print(f"\nAverage Reward over {num_episodes} episodes: {avg_reward}")
```

```
⇒ Episode 1: Total Reward = [9956.234]
Episode 2: Total Reward = [10027.319]
Episode 3: Total Reward = [10112.456]
Episode 4: Total Reward = [10100.748]
Episode 5: Total Reward = [10098.62]
Episode 6: Total Reward = [10084.622]
Episode 7: Total Reward = [10097.509]
Episode 8: Total Reward = [10072.049]
Episode 9: Total Reward = [10129.446]
Episode 10: Total Reward = [10057.532]
```

```
Average Reward over 10 episodes: [10073.653]
```

```

np.random.seed(42)
test_data = pd.DataFrame(np.random.rand(100, 10), columns=[f'feature_{i}' for i
# Use this synthetic data as your test dataset
test_data = (test_data - test_data.min()) / (test_data.max() - test_data.min())


```

```
import matplotlib.pyplot as plt
```

```

# Plot rewards over episodes
plt.plot(rewards)
plt.title("Total Reward per Episode")
plt.xlabel("Episode")
plt.ylabel("Total Reward")
plt.show()

```



```

model = PP0('MlpPolicy', train_env, learning_rate=0.0001, gamma=0.99, n_steps=2
model.learn(total_timesteps=50000)

```



```
| clip_fraction | 0.00137 |
```

clip_range	0.2
entropy_loss	-0.0164
explained_variance	-9.06e-06
learning_rate	0.0001
loss	0.328
n_updates	210
policy_gradient_loss	-0.000612
value_loss	0.824

time/	
fps	157
iterations	23
time_elapsed	299
total_timesteps	47104
train/	
approx_kl	1.6972219e-05
clip_fraction	0.000342
clip_range	0.2
entropy_loss	-0.0131
explained_variance	7.63e-06
learning_rate	0.0001
loss	0.354
n_updates	220
policy_gradient_loss	-0.000242
value_loss	0.846

time/	
fps	157
iterations	24
time_elapsed	311
total_timesteps	49152
train/	
approx_kl	8.688058e-05
clip_fraction	0.0021
clip_range	0.2
entropy_loss	-0.00997
explained_variance	2.92e-06
learning_rate	0.0001
loss	0.464
n_updates	230
policy_gradient_loss	-0.000701
value_loss	1.04

time/	
fps	156
iterations	25
time_elapsed	327
total_timesteps	51200
train/	
approx_kl	1.7708167e-05

clip_fraction	0.000391	
clip_range	0.2	
entropy_loss	-0.0079	
explained variance	-1.29e-05	

```
def random_agent(env, num_episodes=10):
    total_rewards = 0
    for episode in range(num_episodes):
        obs = env.reset()
        episode_reward = 0
        done = False
        while not done:
            # Take random actions in batch form for DummyVecEnv
            action = [env.action_space.sample()] # Wrapping action in a list
            obs, reward, done, _ = env.step(action) # Perform action in batch
            episode_reward += reward
        print(f"Episode {episode + 1}: Total Reward = {episode_reward}")
        total_rewards += episode_reward

    avg_reward = total_rewards / num_episodes
    print(f"\nAverage Reward for Random Agent: {avg_reward}")
    return avg_reward

# Evaluate the random agent
random_agent(test_env)
```

```
↳ Episode 1: Total Reward = [-125.85617]
   Episode 2: Total Reward = [7.6937885]
   Episode 3: Total Reward = [-21.145048]
   Episode 4: Total Reward = [169.25829]
   Episode 5: Total Reward = [-90.26116]
   Episode 6: Total Reward = [-47.93909]
   Episode 7: Total Reward = [41.252026]
   Episode 8: Total Reward = [-37.584785]
   Episode 9: Total Reward = [-4.1457086]
   Episode 10: Total Reward = [53.779324]
```

```
Average Reward for Random Agent: [-5.4948545]
array([-5.4948545], dtype=float32)
```

```
pip install shap
```

```

Requirement already satisfied: shap in /usr/local/lib/python3.10/dist-packa
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/di
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.10/di
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.10/d
Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dis
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pyt
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/di
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/pytho
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p

```

```
# Check the shape of SHAP values
```

```
print(f"SHAP values shape: {shap_values.shape}")
```

```

SHAP values shape: (1, 2587)

```

```
import matplotlib.pyplot as plt
```

```
# Assuming you have the rewards for the trained agent and the random agent
```

```

trained_agent_rewards = [50502.418, 50784.35, 50592.164, 50798.4, 50502.348, 50
random_agent_rewards = [-125.85617, 7.6937885, -21.145048, 169.25829, -90.26116


```


```
# Create the plot
```

```

plt.figure(figsize=(10, 6))
plt.plot(trained_agent_rewards, label="Trained Agent", marker="o")
plt.plot(random_agent_rewards, label="Random Agent", marker="x")
plt.xlabel("Episode")
plt.ylabel("Total Reward")
plt.title("Comparison of Trained Agent vs. Random Agent")
plt.legend()
plt.grid(True)
plt.show()

```

 /usr/local/lib/python3.10/dist-packages/ipykernel/ipkernel.py:283: Deprecat  
and should\_run\_async(code)



Start coding or [generate](#) with AI.

